

## Load data

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

df = pd.read_csv('ecommerce_sales_data.csv')
```

## EDA

```
In [2]: df.isnull().sum()
```

```
Out[2]: TransactionID      0
CustomerID      0
ProductID      0
ProductCategory  0
Quantity      0
Price      0
Discount      0
TransactionDate  0
PaymentMethod   0
CustomerLocation 0
CustomerAge     0
CustomerGender  0
CustomerIncomeGroup 0
CustomerLoyaltyScore 0
dtype: int64
```

```
In [3]: df.info()
df.describe()
```

```
<class 'pandas.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   TransactionID          100000 non-null int64
1   CustomerID             100000 non-null int64
2   ProductID              100000 non-null int64
3   ProductCategory        100000 non-null str
4   Quantity               100000 non-null int64
5   Price                  100000 non-null float64
6   Discount               100000 non-null float64
7   TransactionDate        100000 non-null str
8   PaymentMethod          100000 non-null str
9   CustomerLocation       100000 non-null str
10  CustomerAge            100000 non-null int64
11  CustomerGender         100000 non-null str
12  CustomerIncomeGroup    100000 non-null str
13  CustomerLoyaltyScore   100000 non-null float64
dtypes: float64(3), int64(5), str(6)
memory usage: 16.3 MB
```

Out[3]:

|              | TransactionID | CustomerID    | ProductID     | Quantity      | Price         |               |
|--------------|---------------|---------------|---------------|---------------|---------------|---------------|
| <b>count</b> | 100000.000000 | 100000.000000 | 100000.000000 | 100000.000000 | 100000.000000 | 100000.000000 |
| <b>mean</b>  | 50000.500000  | 5504.000280   | 550.554140    | 2.499750      | 501.469993    |               |
| <b>std</b>   | 28867.657797  | 2600.472561   | 260.256267    | 1.118353      | 287.536905    |               |
| <b>min</b>   | 1.000000      | 1000.000000   | 100.000000    | 1.000000      | 5.010000      |               |
| <b>25%</b>   | 25000.750000  | 3251.000000   | 325.000000    | 1.000000      | 251.537500    |               |
| <b>50%</b>   | 50000.500000  | 5511.000000   | 552.000000    | 3.000000      | 501.580000    |               |
| <b>75%</b>   | 75000.250000  | 7751.000000   | 777.000000    | 3.000000      | 750.397500    |               |
| <b>max</b>   | 100000.000000 | 9999.000000   | 999.000000    | 4.000000      | 1000.000000   |               |

In [4]: `df.drop(columns=['TransactionID', 'CustomerID', 'ProductID'], inplace=True)`In [5]: `df.info()`

```
<class 'pandas.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ProductCategory       100000 non-null  str
1   Quantity              100000 non-null  int64
2   Price                 100000 non-null  float64
3   Discount              100000 non-null  float64
4   TransactionDate       100000 non-null  str
5   PaymentMethod         100000 non-null  str
6   CustomerLocation      100000 non-null  str
7   CustomerAge           100000 non-null  int64
8   CustomerGender        100000 non-null  str
9   CustomerIncomeGroup   100000 non-null  str
10  CustomerLoyaltyScore  100000 non-null  float64
dtypes: float64(3), int64(2), str(6)
memory usage: 14.0 MB
```


In [6]: `df.duplicated().sum()`Out[6]: `np.int64(0)`In [7]: `df.dtypes`

```
Out[7]: ProductCategory      str
Quantity      int64
Price         float64
Discount      float64
TransactionDate      str
PaymentMethod      str
CustomerLocation    str
CustomerAge        int64
CustomerGender     str
CustomerIncomeGroup str
CustomerLoyaltyScore float64
dtype: object
```

```
In [8]: df['TransactionDate'] = pd.to_datetime(df['TransactionDate'])
```

## Graphs and insights

```
In [9]: #heatmap
plt.figure(figsize=(10,6))
sns.heatmap(df.corr(numeric_only=True), annot=True, fmt='.2f', cmap='coolwarm',
            linewidths=0.5, square=True, cbar_kws={"shrink": 0.8})
plt.title('Feature Correlation Matrix', fontsize=14, fontweight='bold')
plt.tight_layout()
plt.show()
```




The correlation analysis indicates that most numerical variables in the dataset have very weak relationships with each other, suggesting independent behavior among customer and transaction features.

```
In [10]: # 2. Product Category Countplot
plt.figure(figsize=(10,5))
sns.countplot(x='ProductCategory', data=df, palette='Set2', edgecolor='black', 1
plt.title('Distribution of Product Categories', fontsize=14, fontweight='bold')
plt.xlabel('Product Category', fontsize=11)
plt.ylabel('Count', fontsize=11)
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
```

```
C:\Users\wardo\AppData\Local\Temp\ipykernel_22360\770576662.py:3: FutureWarning:
Passing `palette` without assigning `hue` is deprecated and will be removed in v
0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effe
ct.

sns.countplot(x='ProductCategory', data=df, palette='Set2', edgecolor='black',
linewidth=0.8)
```



#### Product Category Distribution Insights

The distribution of product categories is almost uniform across all categories.

Clothing, Books, Toys, Electronics, Beauty, Sports, and Home & Kitchen have nearly equal transaction counts.


No single category dominates the dataset, which indicates balanced customer purchasing behavior across categories.

This balanced distribution is useful for unbiased analysis and machine learning modeling.

```
In [11]: # 3. Payment Method Countplot
plt.figure(figsize=(8,5))
sns.countplot(x='PaymentMethod', data=df, palette='pastel', edgecolor='black', 1
plt.title('Payment Method Preference', fontsize=14, fontweight='bold')
plt.xlabel('Payment Method', fontsize=11)
plt.ylabel('Number of Transactions', fontsize=11)
plt.tight_layout()
```

```
C:\Users\wardo\AppData\Local\Temp\ipykernel_22360\1398334644.py:3: FutureWarning:
Passing `palette` without assigning `hue` is deprecated and will be removed in v
0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effe
ct.

sns.countplot(x='PaymentMethod', data=df, palette='pastel', edgecolor='black',
linewidth=0.8)
```



#### Payment Method Preference Insights


Customers are using all payment methods almost equally.

Bank Transfer appears to have slightly higher usage compared to other payment methods.

Cash on Delivery, Credit Card, Debit Card, and PayPal transactions are also consistently high.

The dataset shows diversified payment behavior with no heavy dependency on a single payment method.

```
In [12]: # 4. Customer Age Histogram
plt.figure(figsize=(10,5))
sns.histplot(df['CustomerAge'], bins=20, kde=True, color='teal', alpha=0.7, edge
plt.title('Age Distribution of Customers', fontsize=14, fontweight='bold')
plt.xlabel('Customer Age (years)', fontsize=11)
plt.ylabel('Count', fontsize=11)
plt.axvline(df['CustomerAge'].mean(), color='red', linestyle='--', label=f'Mean:
plt.legend()
plt.tight_layout()
```



#### Customer Age Distribution Insights


Customer ages are distributed fairly evenly between 18 and 70 years.

The average customer age is approximately 43 years.

There is no major age concentration or skewness observed in the dataset.

The business appears to target customers from multiple age groups rather than focusing on a specific demographic.

```
In [13]: # 5. Price Histogram
plt.figure(figsize=(10,5))
sns.histplot(df['Price'], bins=30, kde=True, color='coral', alpha=0.7, edgecolor='black')
plt.title('Price Distribution', fontsize=14, fontweight='bold')
plt.xlabel('Price ($)', fontsize=11)
plt.ylabel('Count', fontsize=11)
plt.axvline(df['Price'].mean(), color='darkblue', linestyle='--', label=f'Mean: {df["Price"].mean():.2f}')
plt.legend()
plt.tight_layout()
```



### Price Distribution Insights


Product prices are spread almost uniformly between low-price and high-price ranges.

The average product price is around \$501.47.

No strong skewness or abnormal clustering is observed in product pricing.

The dataset includes products from multiple pricing segments, making it suitable for comprehensive sales and customer behavior analysis.

```
In [14]: # Product Category vs Average Price
plt.figure(figsize=(10,6))
avg_price = df.groupby('ProductCategory')['Price'].mean().sort_values()
plt.plot(avg_price.values, avg_price.index, 'o-', color='darkorange', linewidth=2)
plt.title('Average Price by Product Category', fontsize=14, fontweight='bold')
plt.xlabel('Average Price ($)', fontsize=11)
plt.ylabel('Product Category', fontsize=11)
plt.grid(True, alpha=0.3, linestyle='--')
plt.tight_layout()
```



### Average Price by Product Category Insights

Sports category has the highest average product price among all categories.

Toys category has the lowest average product price.

Most categories have average prices close to \$500, indicating a balanced pricing structure across the platform.

The small variation in average prices suggests that products are distributed evenly across categories without major pricing bias.

```
In [15]: sns.scatterplot(x='CustomerAge',  
                        y='CustomerLoyaltyScore',  
                        data=df  
                        )  
  
plt.title("Customer Age vs Loyalty Score")  
plt.show()
```



#### Customer Age vs Loyalty Score Insights

There is no strong relationship between customer age and loyalty score.

Customers of all age groups show both high and low loyalty scores.

Loyalty behavior appears to be independent of customer age in this dataset.

This suggests that factors other than age may influence customer loyalty more significantly.

```
In [16]: plt.figure(figsize=(12,6))  
sns.boxplot(x='ProductCategory', y='Price', data=df, palette='pastel', width=0.6,  
            showmeans=True, meanprops={"marker":"D", "markerfacecolor":"red", "markeredg  
plt.title('Price Distribution by Product Category', fontsize=14, fontweight='bol  
plt.xlabel('Product Category', fontsize=11)
```



```
plt.ylabel('Price ($)', fontsize=11)
plt.xticks(rotation=45, ha='right')
plt.grid(axis='y', alpha=0.2)
plt.tight_layout()
```

C:\Users\wardo\AppData\Local\Temp\ipykernel\_22360\2796103620.py:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v 0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='ProductCategory', y='Price', data=df, palette='pastel', width=0.6,
```



### Price Distribution by Product Category Boxplot Insights

All product categories show a very similar price distribution pattern.

Median prices across categories are close to \$500.

The interquartile ranges are also similar, indicating consistent pricing spread across categories.

No category shows extreme outliers or abnormal price behavior.

The dataset appears well-balanced, which is beneficial for comparative analysis and machine learning tasks.

In [17]: `df.head()`

Out[17]:

|   | ProductCategory | Quantity | Price  | Discount | TransactionDate        | PaymentMethod    | Cus |
|---|-----------------|----------|--------|----------|------------------------|------------------|-----|
| 0 | Clothing        | 1        | 801.66 | 0.13     | 2023-12-22<br>10:35:26 | Cash on Delivery |     |
| 1 | Beauty          | 1        | 225.03 | 0.06     | 2024-05-05<br>10:40:33 | Debit Card       |     |
| 2 | Toys            | 2        | 544.99 | 0.24     | 2023-12-04<br>05:22:59 | Debit Card       |     |
| 3 | Beauty          | 2        | 157.02 | 0.01     | 2024-01-02<br>11:11:00 | Bank Transfer    |     |
| 4 | Electronics     | 2        | 155.51 | 0.14     | 2024-06-18<br>01:01:29 | Debit Card       |     |

## Overall Insights

- Customer purchasing activity is evenly distributed across all product categories, indicating balanced market demand without heavy dependency on a single category.
- The platform serves customers from diverse age groups, showing broad customer engagement across different demographics.
- Product pricing remains relatively stable across categories, suggesting a consistent pricing strategy throughout the business.
- Multiple payment methods are actively used by customers, which reflects flexibility and convenience in the purchasing process.
- Customer loyalty does not appear to be strongly influenced by age, indicating that loyalty behavior may depend more on user experience, product quality, or service factors.
- No strong correlation was observed among major numerical variables such as price, quantity, discount, and loyalty score, suggesting that customer purchasing behavior is influenced by multiple independent factors.
- The dataset is highly balanced with minimal bias in categories, pricing, and customer distribution, making it suitable for reliable exploratory analysis and machine learning implementation.
- Overall, the business demonstrates stable transactional behavior, diversified customer interaction, and consistent product engagement across the platform.